San Francisco 311 Requests Data Streaming ETL Pipeline

Project Overview

Summary

San Francisco faces the challenge of efficiently managing and addressing the diverse array of incidents reported through its **311 system**. With a continuous influx of cases, there is a pressing need to **streamline the allocation of resources** and **enhance response times to improve service quality**.

Objectives

1. **Streaming Data ETL Pipeline:** Leveraging skills from Data Streaming course to build a real-time streaming data pipeline.
2. **Data Visualization:** Mapping 311 cases across the city to identify hotspots and patterns.
3. **Closure Time Prediction:** Using machine learning to predict case closure times, considering factors like case type and location.

Architecture Diagram

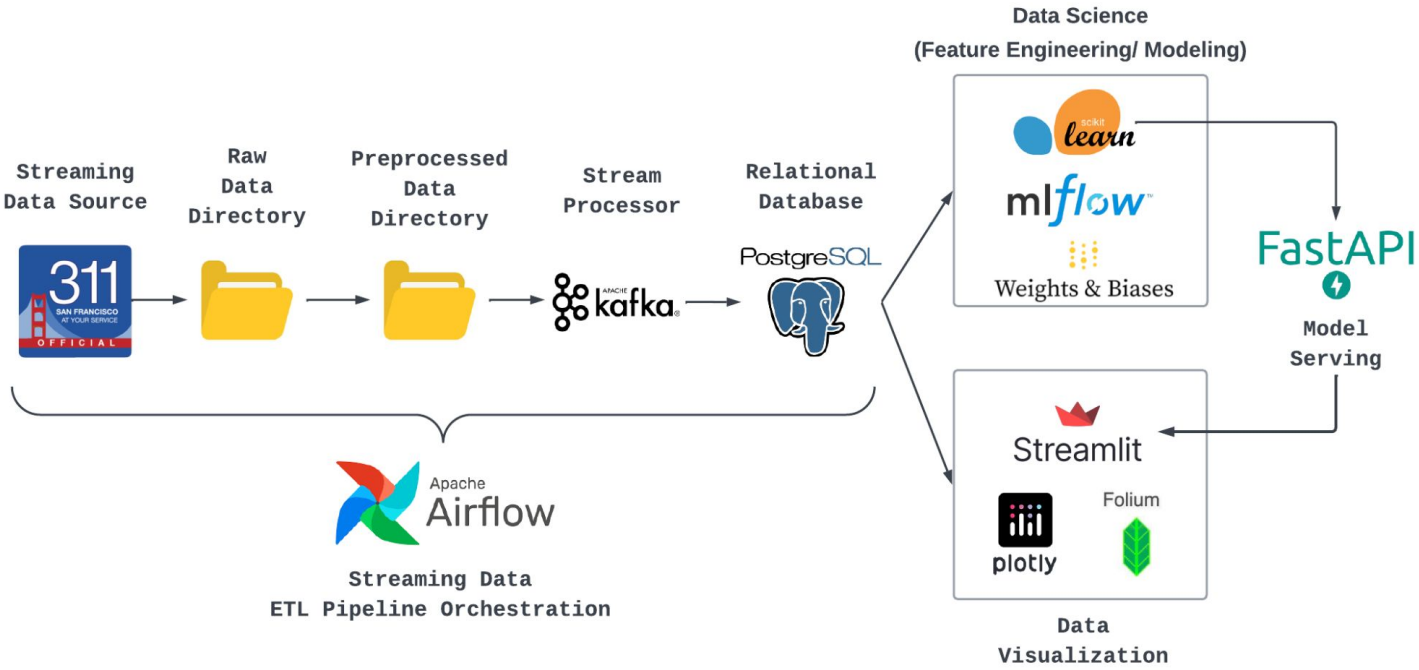

Dataset

SF.GOV

- Open source dataset managed by San Francisco Government
- **Dataset Creation:** October, 2011
- **Update frequency:** Daily (multiple times per hour)
- **Dataset Size:** 2.2 GB
- **Number of Records:** 6.6 M
- **Features:** 48 (interpretable: 15)

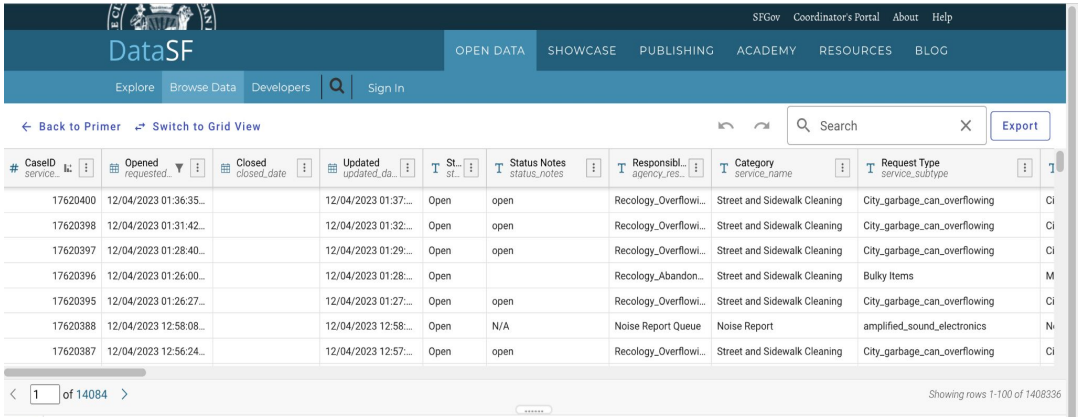

Source: https://data.sfgov.org/City-Infrastructure/311-Cases/vw6y-z8j6/explore

Tools & Packages

- **Data Orchestration:** apache-airflow, sodapy
- **Data Streaming:** avro, confluent_kafka, fastavro
- **Database:** pgAdmin, psycopg2, sqlalchemy
- **Modeling & Serving:** fastapi[all], mlflow, scikit-learn
- **Data Visualization:** geopandas, folium, numpy, pandas, plotly, streamlit
- **Version Control:** git, Github

Airflow - ETL Pipeline Orchestration

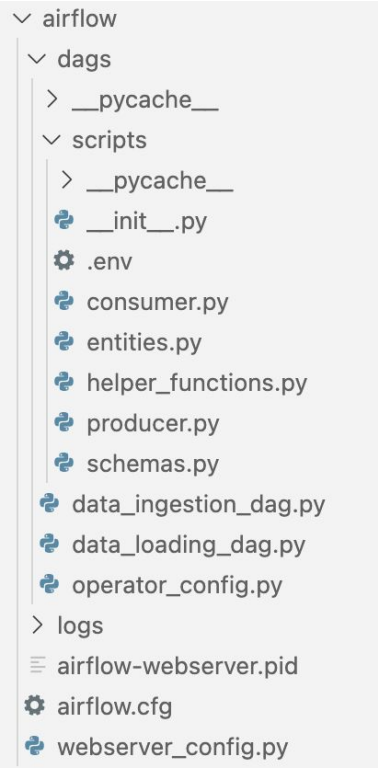

DAG 1

```
with DAG(
    dag_id="sf_311_data_ingestion",
    schedule='15 0 * * *',
    description='Fetch, preprocess, and produce data to Kafka cluster',
    default_args=default_args) as dag:

create_raw_data_dirs_ops >> fetch_data_ops >> save_raw_data_ops
save_raw_data_ops >> extract_cols_ops >> create_processed_data_dirs_ops >> save_processed_data_ops
save_processed_data_ops >> produce_data_ops
```

DAG 2

```
with DAG(
    'sf_311_data_loading',
    default_args=default_args,
    description='Consume data from Kafka and upsert to PostgreSQL Database',
    schedule_interval='30 0 * * *',
):
consume_data_ops
```

Airflow - ETL Pipeline Orchestration

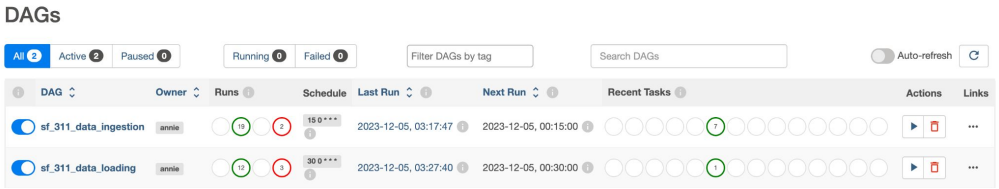

DAG 1 - Data Ingestion, Preprocessing, Produce

Schedule: every night at 00:15 AM

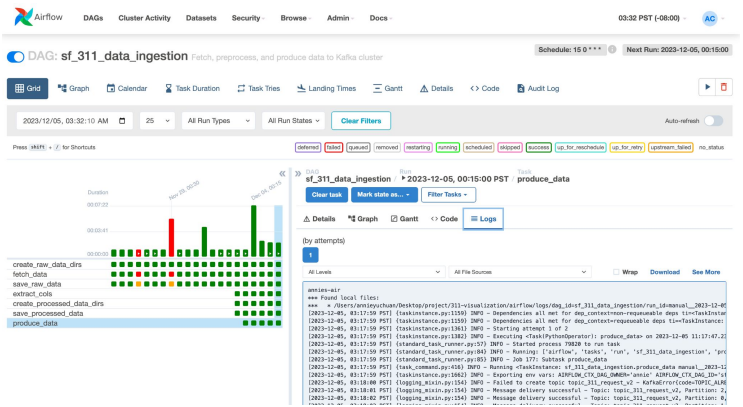

DAG 2 - Consume to Database

Schedule: every night at 00:30 AM

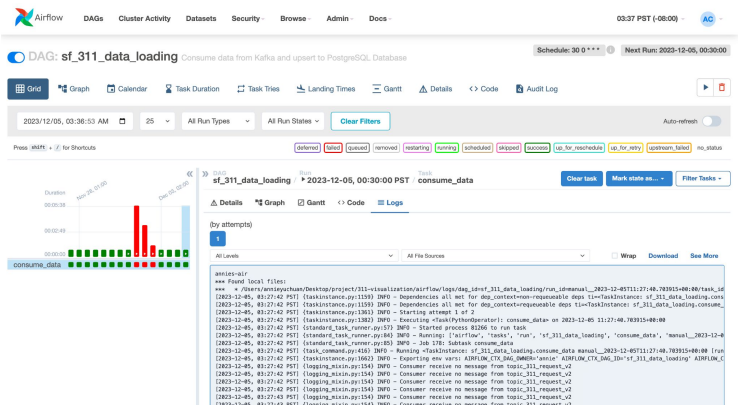

Kafka - Data Streaming

Kafka Topic

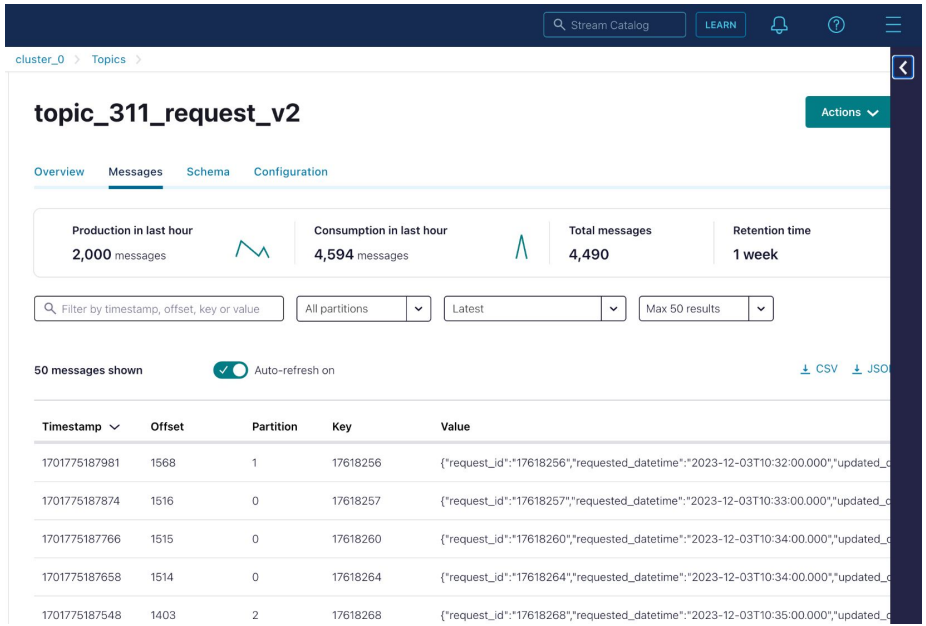

Kafka Schema Registry - Avro Serializer

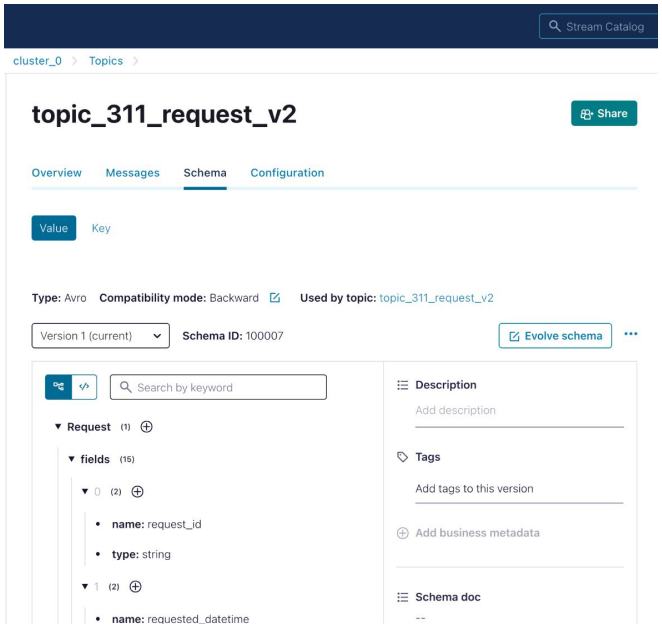

PostgreSQL Database

SQL Table Creation

```
CREATE TABLE IF NOT EXISTS historical_311_request
(
    request_id character varying PRIMARY KEY NOT NULL,
    requested_datetime character varying NOT NULL,
    updated_datetime character varying NOT NULL,
    status_description character varying,
    agency_responsible character varying,
    service_type character varying,
    service_subtype character varying,
    address character varying,
    street character varying,
    supervisor_district character varying,
    neighborhood character varying,
    police_district character varying,
    latitude double precision,
    longitude double precision,
    source character varying
);
```

pgAdmin - PostgreSQL Tool

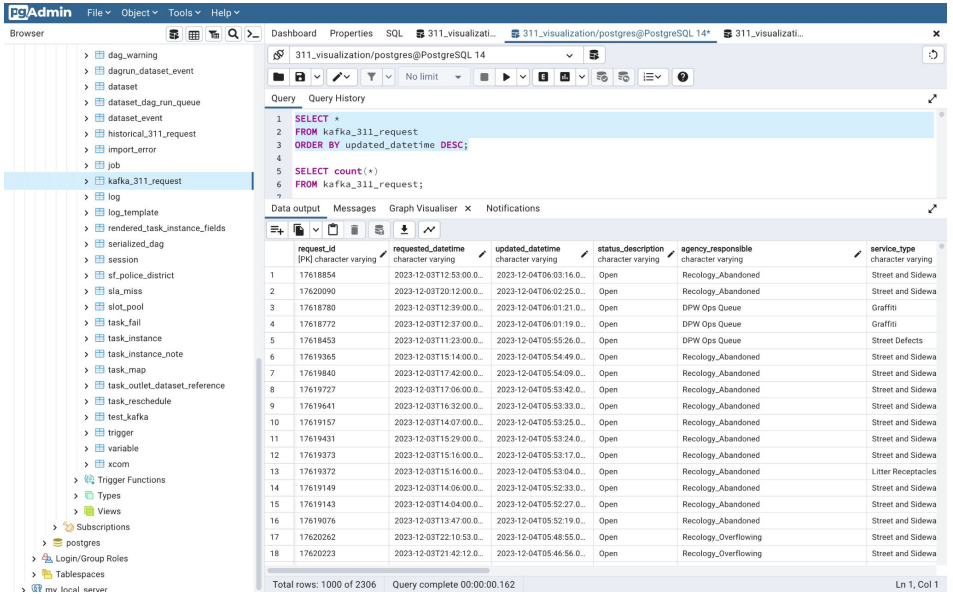

Scikit-Learn/ FastAPI - Modeling & Serving

Scikit learn

```python
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix
import joblib

# Splitting data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Model Training and Hyperparameter Tuning with Random Forest
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10, 15]
}

rf = RandomForestClassifier(random_state=42)
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_

# Training the model with best parameters
best_rf = RandomForestClassifier(**best_params, random_state=42)
best_rf.fit(X_train, y_train)

# Predictions on test set
y_pred = best_rf.predict(X_test)

# Metrics Calculation
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

print(f"Accuracy: {accuracy}")
```

```
Accuracy: 0.586056549770947
```

FastAPI

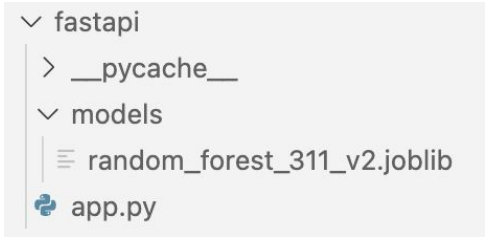

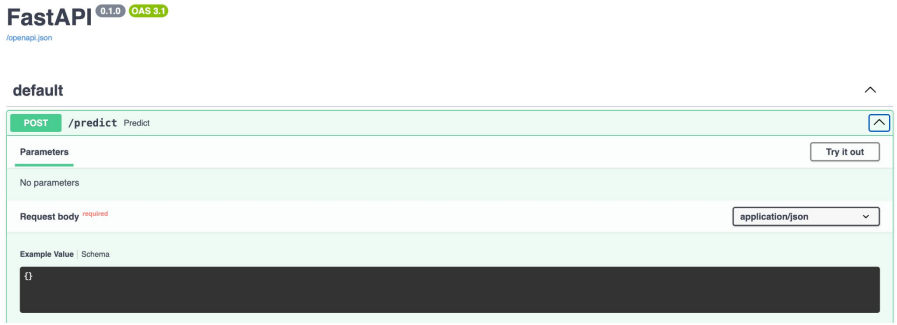

Streamlit/ Plotly/ Folium - Visualization

- streamlit
  - assets
  - helpers
    - __pycache__
    - __init__.py
    - constants.py
    - data_processing.py
    - queries.py
    - utils.py
    - visualizations.py
  - pages
    - .env
    - 1_Dashboard.py
    - 2_Prediciton.py
  - Home.py

Page 1 - Real Time Data Visualization Dashboard

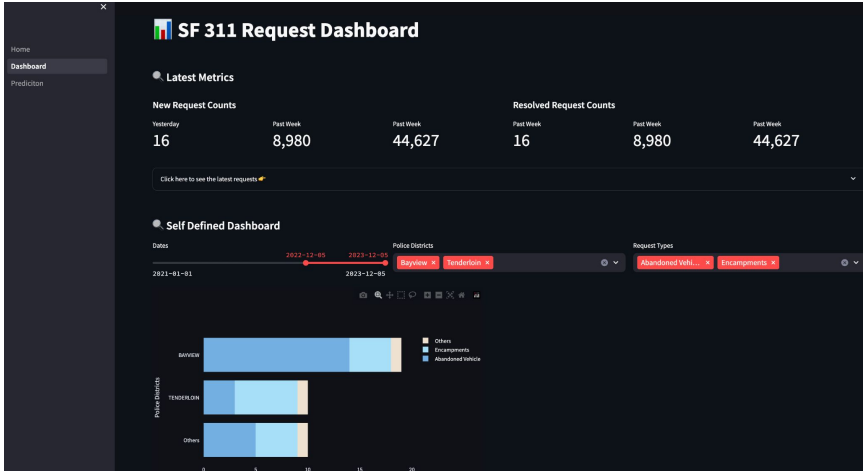

Page 2 - Request Time Prediction

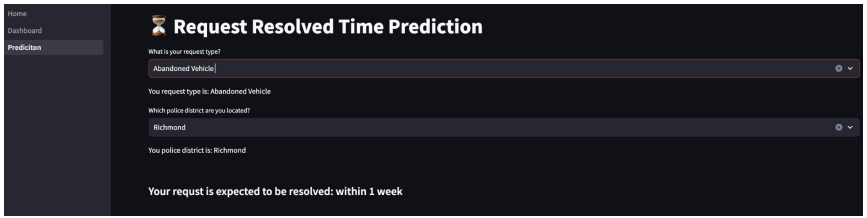

Demo